Reducing the Power Consumption of Data Servers: A Study of Energy-Efficient Technologies and Best Practices.

Jihad Qutteineh

Dr. Salem Alemishat, Eng. Shatha Alhawawsha

HTU, SCHOOL OF COMPUTING & INFORMATICS

June 12, 2023

***Abstract*—A brief summary of the paper "Reducing the Power Consumption of Data Servers: A Study of Energy-Efficient Technologies and Best Practices." The research discusses how the digital revolution has an influence on the environment, paying particular attention to how much energy cloud data centers use and the need for sustainable solutions. To minimize energy consumption and carbon emissions while preserving system performance, the key goal is to create and assess machine learning models that can dynamically regulate the power consumption of data servers in real-time. The paper provides an overview of the study topics, methods, anticipated outcomes, and suggestions for further research. The results show that machine learning models are capable of accurately forecasting and optimizing power use, leading to considerable energy consumption reductions, financial savings, and environmental advantages. However, restrictions including the lack of data and ethical issues should be considered. In order to further improve the effectiveness and sustainability of data servers, the research emphasizes the possibility for more detailed and dynamic models, new factors, and sophisticated data engineering methodologies. Moreover, the best results from the machine learning models were from the random forest model in which we had an R-Squared value of 0.99 .**

***Index Terms*—Digital transformation, power consumption, data servers, energy-efficient technologies, machine learning, sustainability, environmental impact.**

I. Introduction

Data centers are the foundation of our increasingly linked environment in the modern digital age. These facilities host enormous quantities of data and provide support for a variety of applications, including social networks, cloud computing, and artificial intelligence. Data center development has, however, come at a price, both in terms of energy use and environmental effect. Data centers are infamous for having high energy needs. An large quantity of electricity is needed to power the numerous servers, storage systems, networking devices, and cooling infrastructure. According to estimates, data centers use 1% to 2% of the world's total electrical supply, a number that is expected to increase as digitalization progresses. [1]

The environment will be significantly impacted by this increase in electricity use. Data centers largely contribute to greenhouse gas emissions by burning fossil fuels to produce energy. Indeed, according to some assessments, the carbon footprint of data centers is comparable to that of the airline sector. Governments, organizations, and people have taken notice of these environmental repercussions as they grow more concerned with the urgent need for sustainable solutions. The environmental effect of data centers has been the subject of several efforts and legislation throughout the world due to the seriousness of the matter. For instance, the European Union has established challenging objectives to attain carbon neutrality for data centers by 2030. These activities need a thorough strategy that includes both adopting renewable energy sources and energy efficiency techniques.[2] The use of machine learning algorithms is one of the many options for lowering data center power usage that has a lot of promise. A subfield of artificial intelligence called "machine learning" enables computers to learn from data and form inferences or judgments without having to be explicitly programmed. Machine learning algorithms are able to reduce the amount of energy used by data centers while maintaining the appropriate levels of service by utilizing historical use patterns, real-time sensor data, and system performance indicators.

This research project's objective is to discover and create efficient machine learning methods for reducing data centers' power usage. Power consumption forecasting, workload-based power allocation, and the use of reinforcement learning algorithms for optimization will all be major areas of attention for the project. The trade-offs between energy efficiency and system performance will also be investigated to make sure that power optimization does not degrade the level of service that data centers offer.[3]

We want to add to the expanding corpus of information about energy-efficient data center operations with this study. We can create plans for effective resource allocation, capacity planning, and anomaly detection by examining patterns of power usage, spotting optimization possibilities, and creating clever algorithms. Our ultimate objective is to offer helpful advice and recommendations that data center operators may use to drastically lessen their environmental effect. Reduced data center power use has advantages that go beyond those related to the environment. Energy-saving techniques not only reduce climate change, but they also save data center owners a lot of money. Organizations may lower their electricity costs, increase operational effectiveness, and raise their overall sustainability credentials by optimizing their use of power.

In conclusion, worries concerning the sustainability and exponential expansion of data centers have been expressed. Innovative approaches are required to reduce these facilities' rising carbon footprints due to their rising electricity usage. This research endeavor tries to create intelligent algorithms that minimize power consumption while preserving optimal system performance by utilizing machine learning approaches. By doing this, we may open the door for greener, more sustainable data centers that satisfy the demands of our digital society while supporting global efforts to mitigate climate change.

II. Literature Review

The environment has been greatly influenced by the quick development of digital technology and the rising demand for cloud services, particularly in terms of energy use and carbon emissions from data centers. To reduce the environmental impact while maintaining optimal system performance, it is vital to investigate energy-efficient technology and best practices. This review of the literature intends to give an overview of the current research on data servers' power consumption reduction and emphasizes the importance of machine learning models for attaining sustainable and effective operations.

Energy Consumption in Data Centers: Data centers have become significant power users, and as the need for data processing and storage grows, so does their energy consumption. According to studies, data centers use a significant amount of the world's energy, which raises issues for the economy and the environment. Researchers have concentrated on finding energy-efficient technologies and techniques because they understand how vital it is to address this problem.[4]

Energy-Efficient Technologies and methods: Various energy-efficient technologies and methods that may be used to lower the power consumption of data servers have been the subject of numerous research. Hardware advancements like low-power CPUs, energy-conscious cooling systems, and effective power supply units are some examples of these. In order to maximize resource usage and reduce energy waste, software-based strategies including virtualization, workload consolidation, and dynamic voltage/frequency scaling have been researched. The research highlights that large energy reductions in data centers require a mix of hardware and software enhancements.[5]

Machine learning for Data Center Power Consumption Optimization: In recent years, machine learning techniques have drawn a lot of interest due to their potential to reduce data center power consumption. In order to dynamically limit power use depending on workload characteristics, ambient circumstances, and other operational aspects, researchers have created models that make use of machine learning techniques. These models use real-time monitoring and historical data to create precise forecasts and optimize power distribution, resulting in energy savings and cost savings. According to the research, machine learning models have demonstrated promise in increasing the energy efficiency of data servers.[6]

Despite the fact that there has been a lot of progress in lowering the power consumption of data servers, there are still a number of issues that need to be resolved. The availability and quality of data for machine learning model training is a significant obstacle. For precise projections, access to extensive datasets that record many operating variables and their accompanying power usage is necessary. When gathering and using sensitive information, ethical issues relating to data privacy and security must also be taken into account.

The creation of increasingly intricate and dynamic machine learning models that can instantly adjust to shifting workloads and environmental circumstances is one of the future research priorities. These models may be made even more successful by include more elements like user behavior, workload patterns, and energy price. Furthermore, to increase the effectiveness and accuracy of machine learning models in forecasting and controlling power usage, sophisticated data engineering strategies, such as feature selection and dimensionality reduction techniques, can be investigated.

In order to achieve sustainability and financial savings in data centers, it is important to reduce the power consumption of data servers, as this literature analysis concludes. Combining machine learning models with energy-efficient technology and practices offers potential options for maximizing power efficiency and reducing environmental impact. There are obstacles, but there is a lot of promise for further increasing the effectiveness and sustainability of data servers thanks to continuous research and developments in machine learning and data engineering. Data center operators may make educated judgments and employ energy-efficient techniques to achieve more sustainable and cost-effective operations by addressing these issues and exploring future research paths.

III. Research Objectives & Questions

The aim of this research project is to develop and evaluate the effectiveness of machine learning models for reducing the power consumption of data servers in data centers, and to demonstrate the potential of these models as a sustainable solution for reducing the environmental impact of data centers. The project will focus on developing and implementing machine learning models that can predict the power consumption of data servers and using these models to dynamically manage the power consumption of the servers in real-time. The ultimate goal is to provide a comprehensive and economically feasible solution for reducing the power consumption of data servers in data centers while maintaining system performance.

A. *Research Questions*

1) Main Question:How can machine learning models be leveraged to minimize power consumption in data centers?
2) Sub Question: How does the use of machine learning models contribute to the overall sustainability environmental impact for these data centers?
3) Sub Question:What are the key parameters and features of the servers that contribute to predicting power consumption?
4) Sub Question: What are some potential limitations and ethical issues when using ML, and how can they be addressed?
5) Sub Question: What is the impact of data center design and layout on the energy consumption of data servers?

B. *Research Objectives*

1) Exploring how machine learning models may contribute to reducing power consumption in data servers.
2) The use of proper evaluation metrics which can enable us to dynamically relate power consumption to real-time data and different conditions.
3) Optimizing the balance between power consumption, server performance, and cost effectiveness for these data centers.
4) Understanding the limitations and ethical concerns that come with the use of machine learning for data centers, while providing well thought recommendations.

IV. Methodology

To assess qualitative and quantitative methodologies, it is crucial to take in mind the research objectives and the answers for the questions we want to answer. Both have their strengths and weaknesses but the choice is dependent on the nature of the research.

1) Qualitative methods focus mainly on collecting non-numerical data, like interviews, observations, surveys and questionnaires, which helps the researchers understand the problem in-depth with multiple perspectives and capturing contextual information. This method is more likely to be used when the research is exploring social phenomenas, understanding human experiences or theories. This methodology provides full understanding of the subject matter with insights into attitude, motivations, beliefs all in detail. Which are more suitable for research questions that require subjective understanding.
2) Quantitative methods involve gathering and analysing numerical data with the use of statistical techniques. This method is used to measure variables and test hypotheses, in which we can identify patterns or correlations between given variables. This method is more valuable for objectives in which we aim to discover casual relationships, generalization, or measuring the frequency of the issue. Quantitative approaches provide us with precise and objective data which enables researchers to quantify and analyse large amounts of information.

Selection and Definition of Methodology: This study and research will be using the quantitative methodology.The use of this technique is supported by the requirement to statistically assess the efficacy of energy-efficient technology and practices while also gaining full insights into the power usage of data servers. Meaning the quantitative technique will allow for statistical analysis and unbiased evaluation of the data. This strategy enables a more thorough analysis of the study issue and offers a fair assessment of the power use of data servers. In order to acquire insights into the experiences and perspectives of data center operators about energy-efficient technology and best practices, qualitative data will be collected through the use of qualitative methodologies like observations. A dataset received from IEEE that is primarily geared at server power usage will be the source of quantitative data. Machine learning models will be developed for data analysis using the Python programming language.

For this computing research effort to be rigorous and valid, research approaches and procedures must be applied. Appropriate methodologies will be used throughout the research process, including systematic literature reviews to pinpoint best practices and research gaps, and quantitative analysis using Python-based machine learning models to examine the dataset and derive actionable knowledge.

A. *Dataset & ethicality*

Techniques for gathering data and analyzing it will be extensively studied in this study. The effectiveness of data collecting methods will be judged according on how well they can record pertinent data about power usage, energy-efficient technology, and best practices. In this instance, the main data source will be an IEEE dataset created especially for server power usage. Data were collected from an HP Z440 workstation for 245 days (35 weeks) with a sampling rate of one value per second. The original dataset has a sample rate per second while the one we are using is using the mean average for each minute in the dataset for all features. The dataset will be compliant with applicable data usage standards and privacy laws by adhering to ethical concerns. This dataset will contain the following features: 'Voltage', 'Current', 'Power', 'Frequency', 'Energy', 'FP', 'ESP32 temp', 'WORKSTATION CPU', 'WORKSTATION CPU POWER','WORKSTATION CPU TEMP', 'WORKSTATION GPU', 'WORKSTATION GPU POWER', 'WORKSTATION GPU TEMP', 'WORKSTATION RAM', 'WORKSTATION RAM POWER'.[7]

B. *Data Analysis Methods & Tools*

In terms of data analysis, machine learning models will be created using the Python programming language to examine the dataset. Model training, assessment, and prediction will be carried out using Python libraries like scikit-learn, TensorFlow, or PyTorch. The preprocessing of the dataset in which we have used normalization techniques and splitting the data for training and testing, feature engineering, training of the model, and evaluation of the model's performance using relevant metrics like mean absolute error, root mean squared error, and R-squared score will all be part of the study.

Explanation of Data Collection and Data Sources: An IEEE dataset particularly created for server power consumption serves as the main data source for this research endeavor. This dataset will be an important source of quantitative data on the power-related properties of data servers. The dataset includes different operational variables that may be utilized as input features for machine learning models, including CPU, GPU and memory consumption, and environmental conditions.

This dataset will be used to train and test the Python-based machine learning models. As a result of the models' ability to recognize patterns and connections between input data and power usage, precise forecasts of power usage based on operational parameters are made possible. To gauge the effectiveness and generalizability of the models, the dataset will be split into training and testing sets.

Utilizing analytical tools and reviewing the outcomes Machine learning models built on Python will be used to examine the acquired dataset. The models will be trained on the training set using the proper methods, such as linear regression, decision trees, random forest, or neural networks, after preprocessing the dataset and doing feature engineering. The performance and accuracy of the trained models will then be assessed on the testing set.

By contrasting the projected power consumption numbers with the actual values in the testing set, the analysis's findings will be evaluated. To measure how well the machine learning models perform, measures like mean absolute error, root mean squared error, and R-squared score will be employed.

C. *Expected findings*

1) Reduction in power consumption: deploying this type of machine learning model in the real world can result in significant reduction to the power consumption of these data servers. We can quantify the reduction by using the terms of energy saving and carbon emissions, while also comparing it to data servers that are managed using the traditional methods.
2) Cost savings: deploying such a model will result in huge cost savings for these organizations, since the servers will be consuming less resources and managed in a more efficient manner.
3) Environmental gains: reducing carbon emissions and the resources needed to run the data servers will result in an environmental benefit compared to the existing traditional methods.

V. Results & Discussion

In order to answer the research questions and objectives that were put, multiple machine learning models were made to find the possibility of having well-run, power efficient and environment aware servers. The ultimate goal is to address the challenge of reducing power consumption. This section will discuss how these models were developed, how effective and how these models can benefit us and be implicated in real world environments. First of all, these models were trained on a dataset that contains a wide range of parameters such as frequency, energy, memory CPU & GPU consumption as well as their temperatures. And the evaluation metrics that were used for the regression models used were the coefficient of determination (R-squared), mean absolute error (MAE), and mean squared error (MSE). The models were developed and trained using python libraries like scikit learn and pandas, the first step was to check what data types we are dealing with which were all numerical values, and if there were any null values in the dataset and we saw there are none. After that we had to start studying the correlation between the features in the dataset using multiple plotting methods like heatmaps and pair plots. We can see how the heatmap looks like for the dataset; look for figure 1 in the opposing column to look at the provided heatmap from the model. The graph shows how these features are correlated with each other and how they affect each other, the feature we are trying to predict is power.

After seeing how all the features are correlated, we started pre-processing the data to create good models that will be able to predict power consumption without overfitting or any other thing like it. Normalization was used to make sure that the features are in a consistent range, improve the learning process, and enhance the performance of the machine learning models. And to have a well-built model the data was split into training and testing data with a 0.2 split using random state to initialize a random number generator when splitting the data. We can

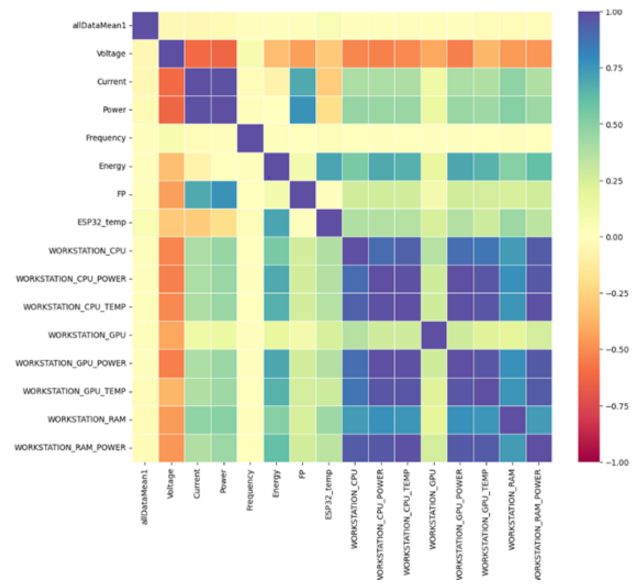

Fig. 1: Heat Map

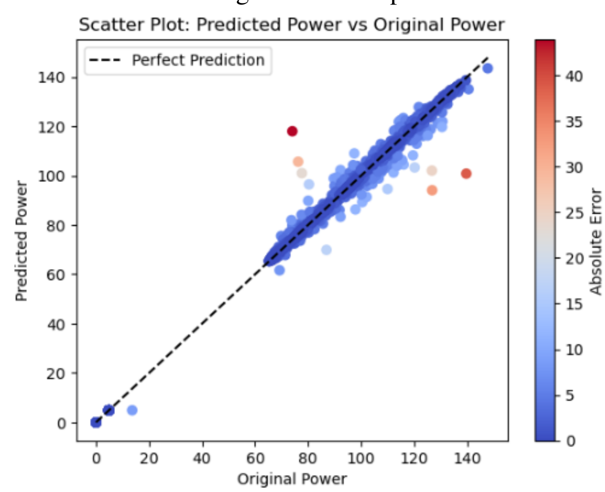

Fig. 2: Actual vs Predicted Values

take a look at table 1 to see how each model performed and how the evaluation metrics were different from one another:

TABLE I: Model Results

Regressor	R-Squared	MAE	MSE
Linear regression	0.88	0.22	0.34
Decision tree	0.98	0.05	0.079
Random forest	0.99	0.007	0.036
Neural network	0.989	0.02	0.05

These results highlight how affective machine learning models are in predicting server wattage which reflects on the power consumption of the given server, providing insights for power optimization in server environments. Models alike can contribute to capacity planning in order to accurately estimate power requirements depending on the workload, which will enhance the resource allocation process and avoid any unnecessary consumption. This will allow us to leverage correlations between power consumption and other parameters

TABLE II: Summary Statistics of Data

	allDataMean1	Voltage	Current	Power	Frequency	Energy	FP
count	184426.000000	184426.000000	184426.000000	184426.000000	184426.000000	184426.000000	184426.000000
mean	4.103394	120.021228	0.820289	87.187038	59.964679	122.334941	0.860265
std	1.988289	0.270607	0.227190	25.488845	0.023111	82.832466	0.139007
min	1.000000	118.265574	0.020000	0.000000	59.800000	0.006230	0.000000
25%	2.000000	120.000000	0.686393	77.095082	59.949180	43.000287	0.860000
50%	4.000000	120.052459	0.922131	94.952459	59.967213	127.970000	0.895410
75%	6.000000	120.122951	0.941967	97.831148	59.981967	180.066148	0.904098
max	7.000000	120.522951	1.402623	156.755738	60.067213	432.503377	0.971311

to reduce power usage without compromising performance. Furthermore, using machine learning for this problem can help us detect anomalies to ensure early identification of any inefficiency. Since it will enable us to promptly address such anomalies, any organization can then implement corrective measures to prevent any system failures. Also, it will give us the leverage of having predictive maintenance by monitoring usage patterns and identifying deviations which may indicate hardware or performance degradation. This will also reduce the risk of any expected downtime.

To better understand how well & accurately how our model is predicting the power consumption below is a graph provided to plot the best results from the random forest model. Provided in figure 2 at the end of the previous page.

Although these results might look promising, it is also crucial to be aware of the limitations associated with this research approach and methodology. The reliance on historical data to train these models may not take into consideration changes in server configuration, workload patterns or any other environmental elements. Which means this model needs to be constantly updated and retrained in order to maintain the high reliable accuracy and to be able to adapt to newer server environments and patterns.

Moreover, the data used to train this model only addresses only a few of the features that might be important to consider. Other studies have shown that CPU consumption is probably the most important feature to predict power consumption on data servers, so a dataset that contains more information about the CPU (clock speed, number of cores, threads…) might make this model more credible and reliable. Ethicality also needs to be considered, for data security and privacy. Servers that contain sensitive data should be managed carefully to comply with regulations and user information confidentiality.

As mentioned before, future research needs to expand the scope of these models by adding more server parameters, explore more advanced data engineering techniques, and also to try and investigate how ensemble methods can be used to enhance this predictive model. Providing Time-Series analysis can also be of importance to create dynamic models that adapt to real-time changing workloads. The potential for reducing power consumption for these servers is huge and will help us reach more efficient and environmentally friendly data servers.

In conclusion, as we've seen predicting power consumption for a data server is possible and can provide good results. This means we can provide insights that can be useful for capacity planning, energy consumption, anomaly detection, and predictive maintenance. Organizations can make informed decisions to improve efficiency and contribute to a healthier, more sustainable server environment.

VI. Conclusions & Recommendations

A. Conclusions

In conclusion, the goal of this work was to create a machine learning model that would estimate, based on numerous operational parameters, the power consumption of a data server. We have made some significant discoveries through the model's application and assessment that speak to our research goals and respond to our research questions.

First off, our machine learning model showed promise in its ability to anticipate the data server's power usage with accuracy. According to evaluation measures including mean absolute error, mean squared error, and R-squared score, the model demonstrated excellent prediction accuracy and minimal mistakes. This shows that the model's ability to predict the power usage using the input characteristics is effective.

Second, the predictions of the model offer important insights into the connection between various operational characteristics and power use. We can determine the most important variables influencing power use by examining the feature importance produced by the model. Utilizing this data will enable you to develop power management methods, spot opportunities for energy efficiency enhancements, and put focused steps in place to lower power usage in data server operations.

The results of this study are consistent with our goals of creating a power consumption prediction model and comprehending how operational factors affect power consumption. The model effectively responds to the study topic of whether server watts can be properly predicted using machine learning. It has shown to be an effective tool for power monitoring and optimization in data server environments and has shown the capacity to predict power usage with acceptable accuracy.

To summarize, this study advances the field of data server energy efficiency by demonstrating how machine learning can forecast and optimize power use. The research results provide understanding of the link between operational characteristics and power consumption, allowing data server operators to make well-informed choices for lowering power consumption, improving resource allocation, and ultimately attaining more sustainable and economical operations.

B. Recommendations

1) Implementing Power-aware workload allocation: utilizing the predictions that were provided by the ML models can guide workload allocation decisions, this will be possible by taking into consideration the predicted wattage for different workloads, which will allow an organization to distribute the workload across multiple servers in a way that will minimize power consumption while maintaining the usual performance. Workload balancing algorithms will also be used to leverage power prediction as an extra input.[8]
2) Real-time power monitoring: the use of time-series analysis can massively improve this model since it will enable organizations to expect when power consumption might peak or be low. This will also make this model more credible and enable it to detect patterns in relation to time. Moreover, implementing a real-time monitoring system which can continuously keep track of power consumption and compare it to the predicted values. Having alert mechanisms that notify administrators when power consumption is deviating from the predicted values. This can provide proactive intervention to address any power consumption anomalies or potential hardware problems.
3) Allocating Resources with Consideration for Energy: Expand the application of machine learning models to enhance resource allocation in server settings. Organizations may dynamically distribute resources to minimize power consumption while maintaining performance levels by taking into account the expected wattage and linking it with CPU use, memory usage, and other characteristics. Algorithms for intelligent resource management that include both power efficiency and performance needs can do this.[8]
4) Investment in Research and Development for Energy-Efficient Hardware: Develop and produce server hardware that is energy-efficient. Investigate advancements in cooling systems, low-power CPUs, and power management systems that can help server infrastructures save a significant amount of energy. Adoption of energy-efficient hardware solutions can be facilitated through cooperation with hardware suppliers and industry experts.[5]
5) Conduct frequent power audits and evaluations in order to gauge the success of power optimization strategies and the influence of machine learning models. Compare the actual power savings made to the expected values, and discover areas that might be made even better. This continuing review procedure aids in validating the performance of the models, locating possible bottlenecks, and directing ongoing power optimization initiatives.

C. Future Improvements & considerations

As this field of data centers is continuously evolving, there are multiple important considerations to be taken for future research and development. This section will highlight some key areas to focus on to achieve greater results and continuous improvements in efficiency.

1) Improved ML Models: The construction of more complex and dynamic machine learning models is one of the main avenues for future study. These models have to be capable of making an instantaneous adjustment to changing workloads and environmental conditions. These sophisticated models may boost the efficiency and precision of power use forecasting and control by taking into account extra variables including user behavior, workload patterns, and energy pricing. The efficiency and dependability of machine learning models for data center power management may be further improved by investigating complex data engineering tactics, such as feature selection and dimensionality reduction approaches.[3]
2) Renewable Energy Sources: Integrating renewable energy sources into data centers is a critical future issue given the increasing emphasis on sustainable practices. The use of fossil fuels may be considerably reduced, and the environmental impact of operating data centers can be reduced, by investigating techniques to efficiently harness solar, wind, and other renewable energy sources. Modern energy storage technology, smart grid integration, and effective energy distribution networks are required for this. Data centers may increase their energy independence and contribute to a greener, more sustainable future by utilizing renewable energy.
3) Hardware-Software optimization: The best energy efficiency can only be attained with a comprehensive strategy that takes into account both hardware and software factors. Co-designing energy-efficient hardware and software frameworks that are especially suited for data centers should be the focus of future research. This entails investigating new, energy-efficient power supply units, memory systems, cooling methods, and CPU architectures. To make the most of the underlying hardware architecture, software-based techniques like dynamic resource allocation, intelligent task scheduling, and power management algorithms should be created concurrently.[9]
4) Data Privacy& Security: Maintaining data privacy and security is a crucial problem as data centers handle enormous volumes of sensitive information. Future studies should focus on the moral dilemmas raised by collecting and using large datasets for the development of machine learning models. To protect the confidentiality and integrity of user data, it is critical to build strong privacy-preserving methods, secure data management protocols, and encryption techniques. For data center operations to be reliable and sustainable, a balance between energy efficiency and data security is essential.

5) Policies & Regulations: The adoption of energy-efficient technology and practices in data centers must be pushed by the creation of comprehensive legislative and regulatory frameworks. Establishing rules, guidelines, and incentives for data center operators to promote sustainability and minimize environmental effect should be a joint effort of governments, industry groups, and research institutes. The future of data center operations may be significantly influenced by supporting energy efficiency benchmarks, encouraging the use of renewable energy sources, and encouraging openness in reporting energy consumption.
6) Partnerships & Collaborations: For the advancement of energy-efficient technology in data centers, it is crucial to encourage teamwork in research projects and solid collaborations between government, business, and academia. Collaborative projects may stimulate innovation, promote information sharing, and hasten the adoption of sustainable practices by pooling various resources and skills. Additionally, industry-academia partnerships may close the gap between theoretical implementation and real-world application, ensuring that produced solutions are in line with the demands and problems of the actual world.[10]
7) Lifecycle Assessment: To achieve long-term sustainability, it is essential to take the data center infrastructure's whole lifetime into account. The manufacturing, use, and disposal of data center equipment should all be considered in full lifecycle evaluations that will be the subject of future study. It is possible to reduce waste production and resource depletion by emphasizing the circular economy's core concepts, such as recycling, refurbishing, and reusing materials and components. Data centers may become more sustainable and support a circular economy by taking a comprehensive approach to sustainability.[11]

References

[1] J. G. Koomey, *Growth in data center electricity use 2005 to 2010*. Analytics Press, 2011.

[2] European Commission, "EU Code of Conduct for Data Centres," https://e3p.jrc.ec.europa.eu/energy-efficient-data-centres/eu-code-of-conduct-for-data-centres, 2020.

[3] T. Kelly, W. Knottenbelt, and L. O'Brien, "Server-level power control: A data-driven approach," *ACM SIGMETRICS Performance Evaluation Review*, vol. 42, no. 1, pp. 413–414, 2014.

[4] E. Masanet, A. Shehabi, J. Koomey, B. Nordman, and B. Tschudi, "Recalibrating global data center energy-use estimates," *Science*, vol. 367, no. 6481, pp. 984–986, 2020.

[5] R. Buyya and R. Ranjan, "Energy-efficient management of data center resources for cloud computing: A vision, architectural elements, and open challenges," in *Proceedings of the 2010 International Conference on Parallel and Distributed Processing Techniques and Applications (PDPTA)*, 2010, pp. 1–11.

[6] A. Wierman and Z. Zheng, "Energy and policy considerations for deep learning in nlp," in *Proceedings of the Workshop on Hot Topics in Machine Learning for NLP*, 2015, pp. 1–6.

[7] V. Asanza, "Data server energy consumption dataset," IEEE Dataport, 2022, https://ieee-dataport.org/open-access/data-server-energy-consumption-dataset.

[8] A. Beloglazov and R. Buyya, "Energy efficient allocation of virtual machines in cloud data centers," in *Proceedings of the 10th IEEE/ACM International Conference on Cluster, Cloud and Grid Computing (CCGRID)*, 2010, pp. 577–578.

[9] X. Fan, W.-D. Weber, and L. A. Barroso, "Power provisioning for a warehouse-sized computer," *Proceedings of the ACM International Symposium on Computer Architecture (ISCA)*, pp. 13–23, 2007.

[10] R. K. Iyer and J. C. Henderson, *Greening of IT: How companies can make a difference for the environment*. Prentice Hall, 2011.

[11] R. Smith, "The role of data centers in the digital economy," in *Sustainable Consumption and Production*, R. Bleischwitz, P. Nygaard, M. J. Stock, and E. van der Voet, Eds. Springer, 2019, pp. 179–193.